



Women in ML

# Keras: Shortcut to AI Mastery



Divyashree Sreepathihalli (she/her)  
Software Engineer, Google  
@divyashreess



# Introduction

Software Engineer

Keras Maintainer

KerasCV Lead and Key Maintainer

Specialized in Keras Modelling APIs  
and KerasCV



# Setup

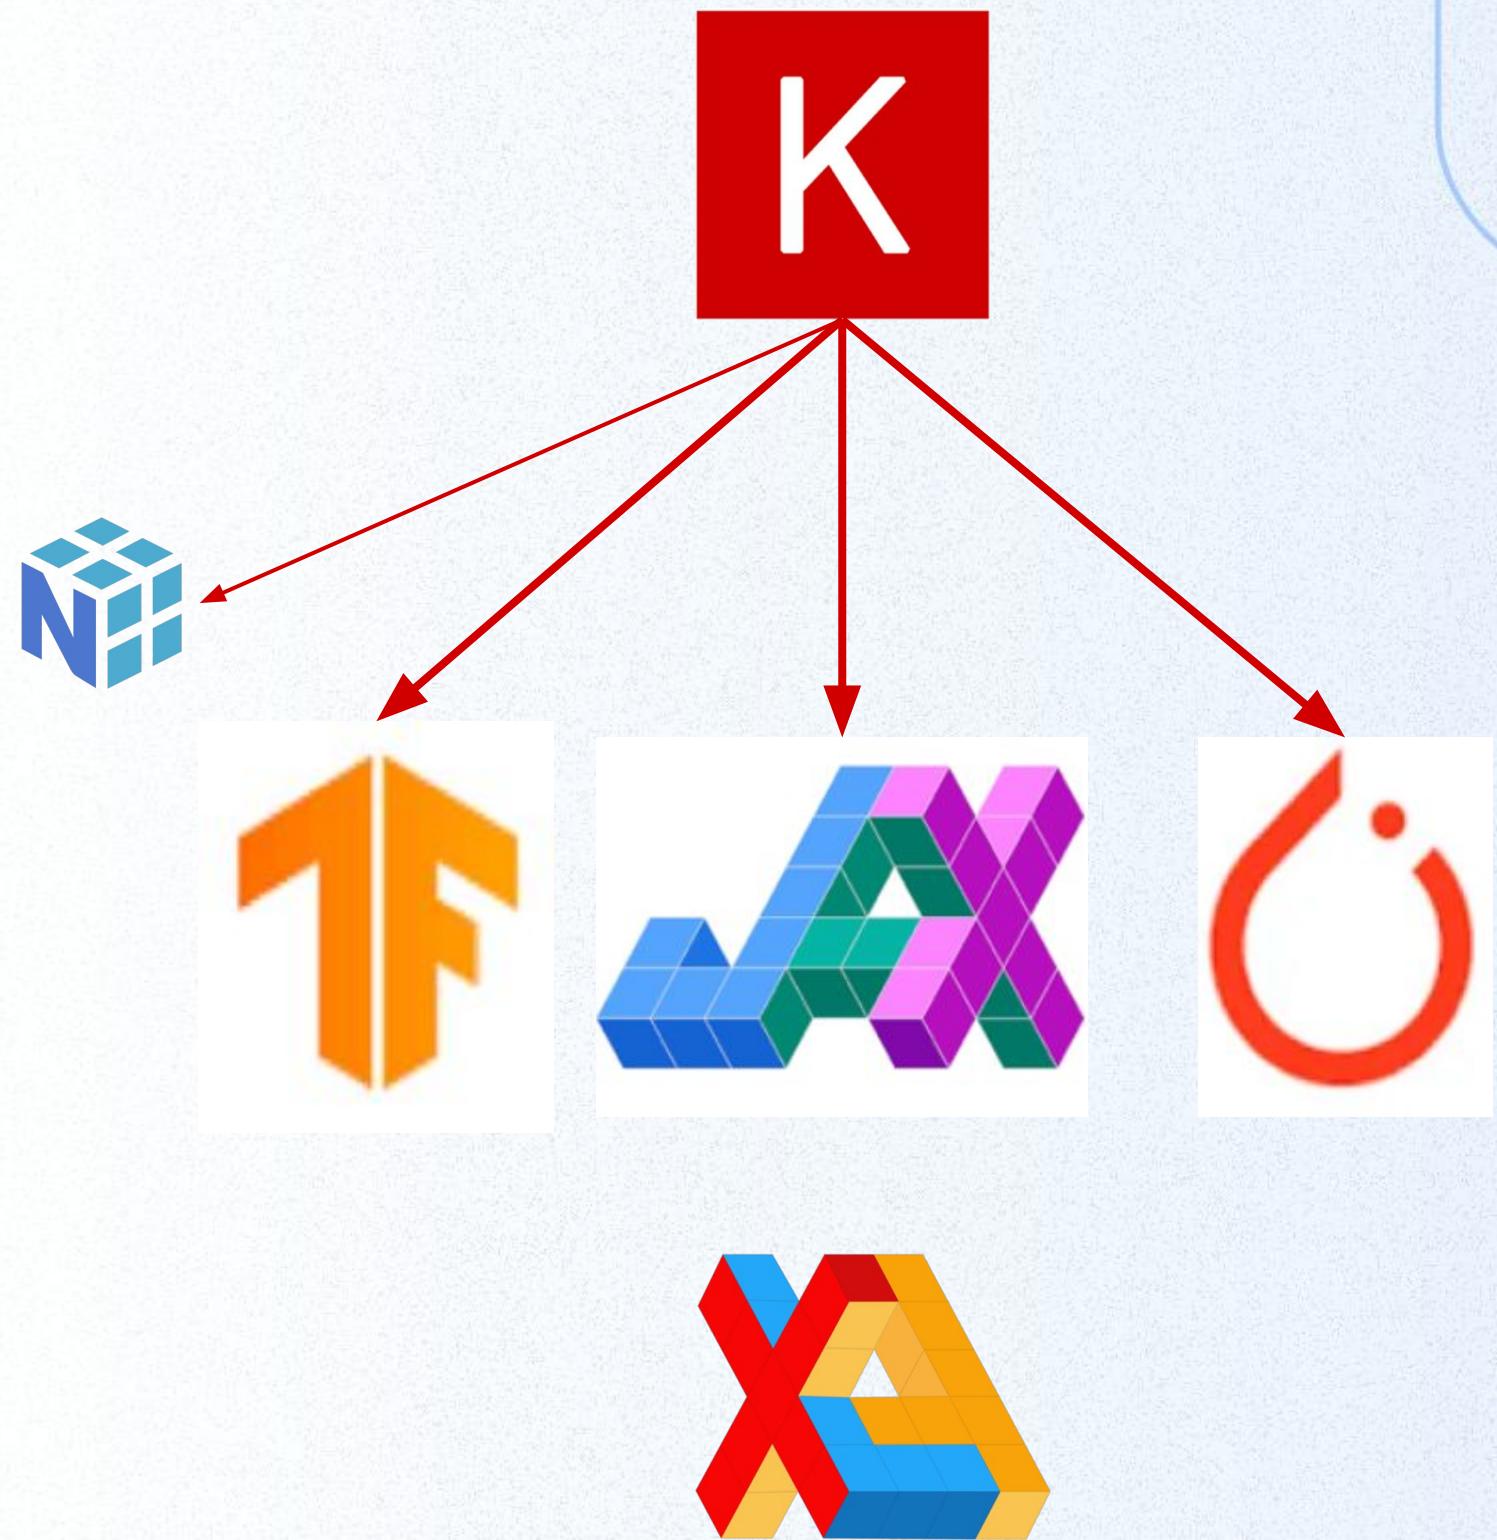
Google Colab Signup - <https://colab.research.google.com/>

Github Repo for todays workshop -  
[https://github.com/divyashreepathihalli/WIML\\_2023](https://github.com/divyashreepathihalli/WIML_2023)



# Multi-backend Keras is back

1. Full rewrite of Keras  
Now only 45k loc instead of 135k
2. Support for TensorFlow, JAX, PyTorch, NumPy  
backends(NumPy backend is inference-only)
3. Drop-in replacement for `tf.keras` when using  
TensorFlow backend
4. XLA compilation by default in TF and JAX



# In today's workshop you will

1. Explore multi backend Keras 3.0 - how to write a simple model, and run it with a backend of your choice, including PyTorch, TensorFlow, or JAX
2. Learn deep learning best practices
3. Learn about the KerasNLP and KerasCV modules and how to build powerful AI applications



BEGINNER



INTERMEDIATE



ADVANCED



EXPERT

# Keras mission

Our goal is to ease the process of building and training your model by offering consistent & simple APIs



# Keras Ecosystem

## Modelling API

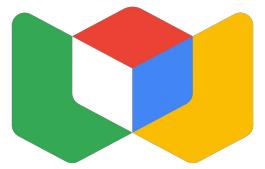
- Keras layers
- Functional/Sequential API
- Keras model interface
- Keras applications

## Training API

- `model.fit()`,  
`model.predict()`
- Loss functions, metrics
- Optimizers
- Callbacks
- Preprocessing layers

## Domain Packages

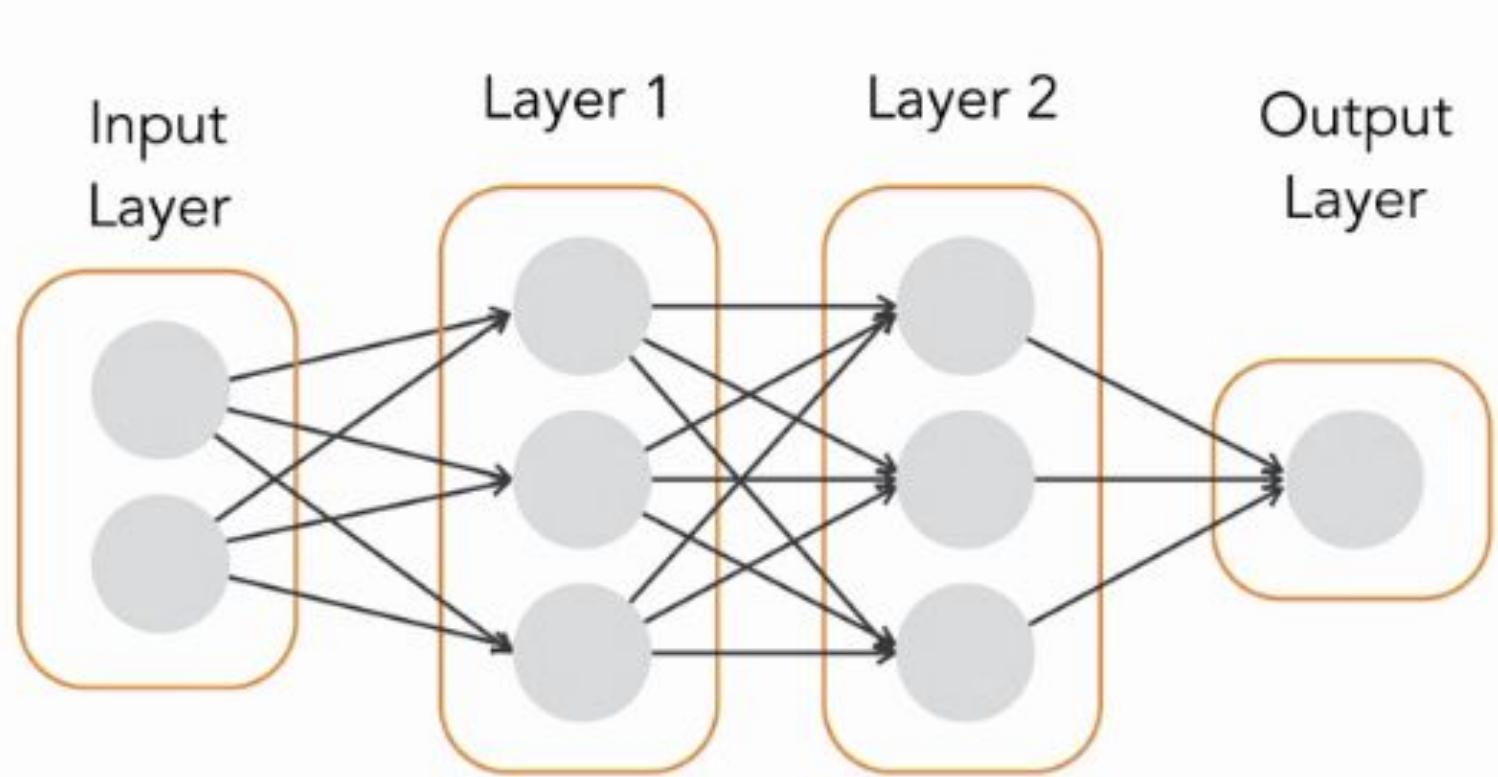
- KerasCV
- KerasNLP



# Keras Modelling APIs



# Keras layers API



```
from tensorflow.keras import layers  
layer = layers.Dense(32,  
activation='relu')  
inputs = tf.random.uniform(shape=(10,  
20))  
outputs = layer(inputs)
```

# Keras Models API

## Sequential

which is very straightforward (a simple list of layers), but is limited to single-input, single-output stacks of layers (as the name gives away)

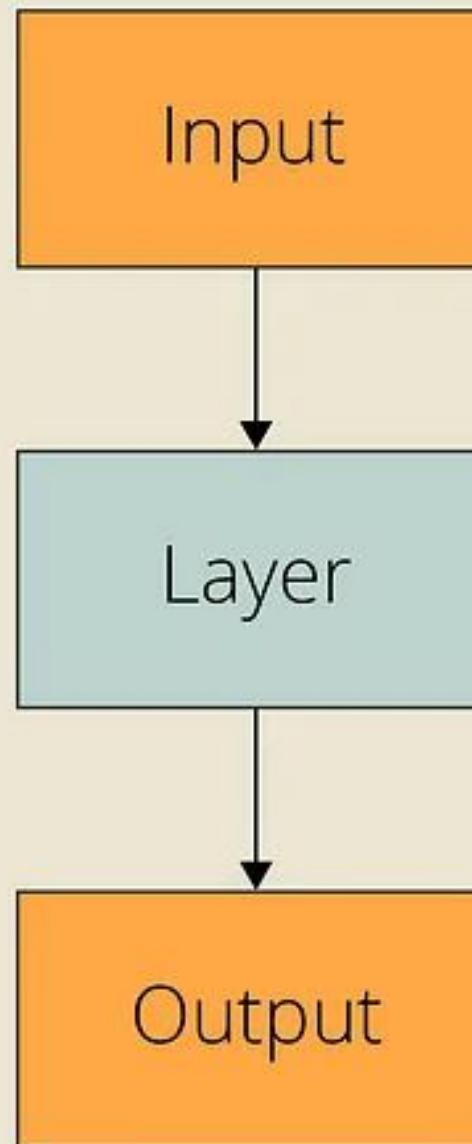
## Functional

an easy-to-use, fully-featured API that supports arbitrary model architectures. For most people and most use cases, this is what you should be using. This is the Keras "industry strength" model.

## Subclass Model

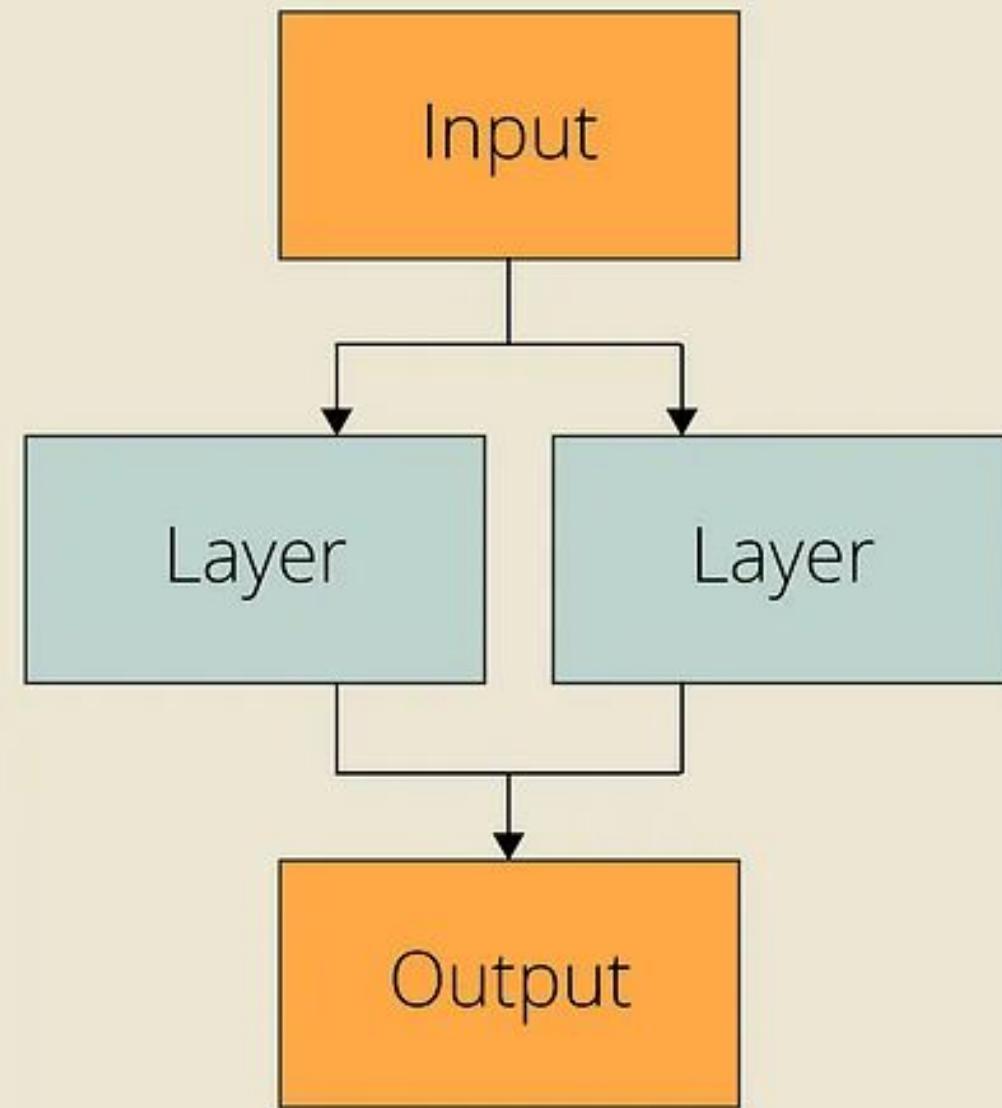
where you implement everything from scratch on your own. Use this if you have complex, out-of-the-box research use cases.

## Sequential API



```
model = tf.keras.Sequential()  
model.add(tf.keras.Input(shape=(16,)))  
model.add(tf.keras.layers.Dense(8))  
model.summary()
```

## Functional API



```
inputs = keras.Input([28, 28, 3])
x = keras.layers.Conv2D(8, 2)(inputs)
x = keras.layers.MaxPool2D(2)(x)
x = keras.layers.Flatten()(x)
x = keras.layers.Dense(2)(x)
outputs = keras.layers.Softmax()(x)
```

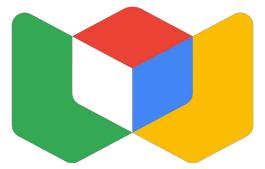
```
functional_model =
keras.Model(inputs=inputs,
outputs=outputs)
functional_model.summary()
```

## Model Subclassing

```
tf.keras.Model
```

```
def __init__():  
...  
def call():  
...
```

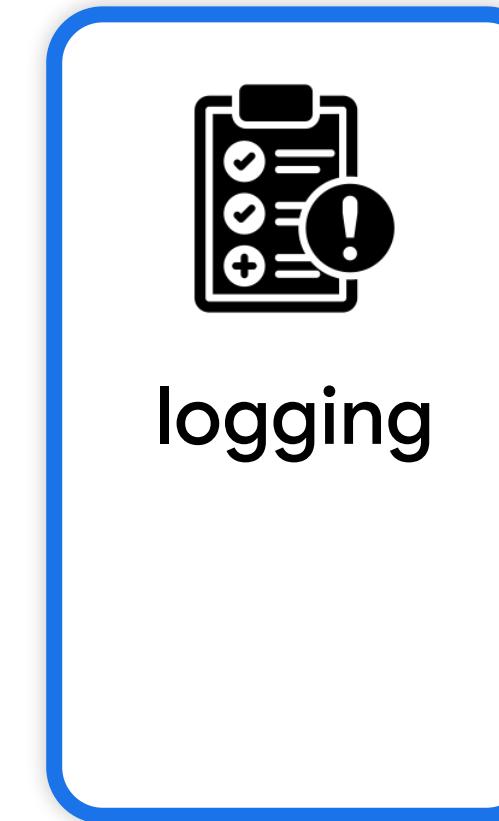
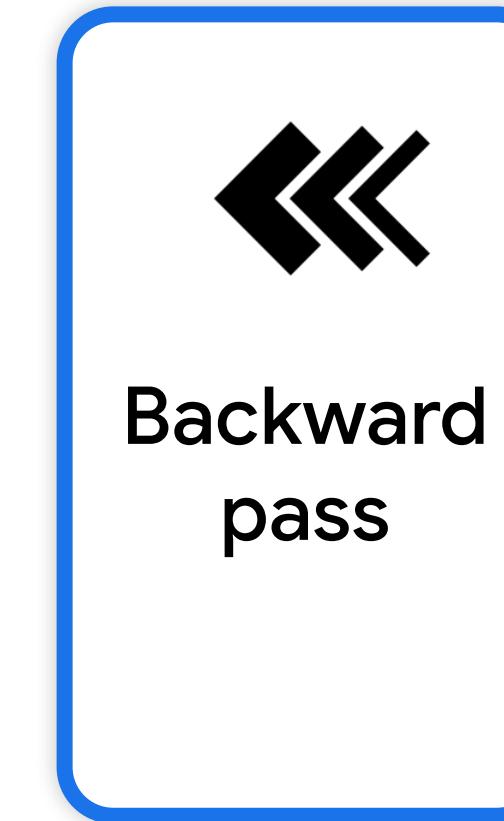
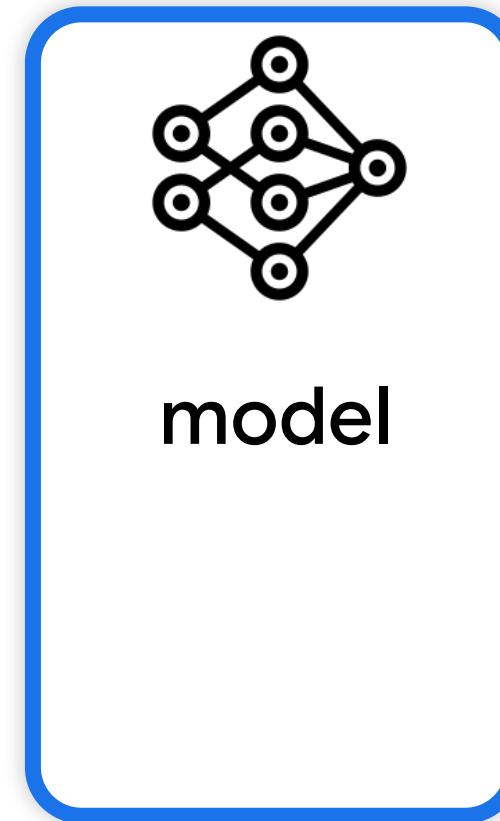
```
class MyModel(keras.Model):  
    def __init__(self, *args, **kwargs):  
        super().__init__(*args, **kwargs)  
        self._dense = keras.layers.Dense(2)  
        self._softmax = keras.layers.Softmax()  
  
    def call(self, inputs):  
        x = self._dense(x)  
        return self._softmax(x)  
  
subclass_model = MyModel()  
  
subclass_model.summary()
```



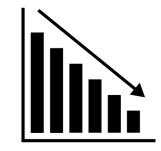
# Keras Training APIs



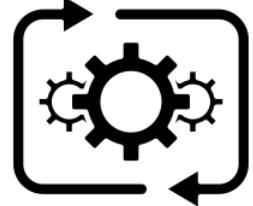
# Training workflow



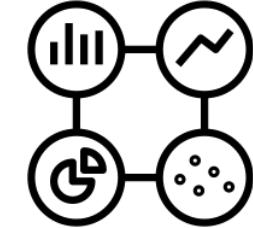
# Keras Training components



Loss



Optimizers



Metrics

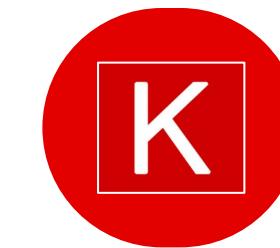


callbacks

```
Model.compile(  
    optimizer="rmsprop",  
    loss=None,  
    metrics=None,  
    loss_weights=None,  
    weighted_metrics=None,  
    run_eagerly=None,  
    steps_per_execution=None,  
    jit_compile=None,  
    pss_evaluation_shards=0,  
    **kwargs  
)
```

```
Model.fit(  
    x=None,  
    y=None,  
    batch_size=None,  
    epochs=1,  
    callbacks=None,  
    validation_data=None,  
    steps_per_epoch=None,  
    validation_steps=None,
```

# Develop cross-framework components with keras.ops



`ops.matmul`



`tf.matmul`



`jax.numpy.matmul_x`



`torch.matmul`



`np.matmul`



Includes the NumPy API – same functions, same arguments

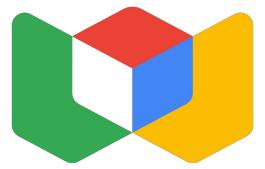
- `ops.matmul`, `ops.sum`, `ops.stack`,  
`ops.einsum`, etc.

Plus neural network-specific functions absent from NumPy

- `ops.softmax`,  
`ops.binary_crossentropy`, `ops.conv`, etc.

Models / layers / losses / metrics / optimizers written with Keras APIs work the same with any framework

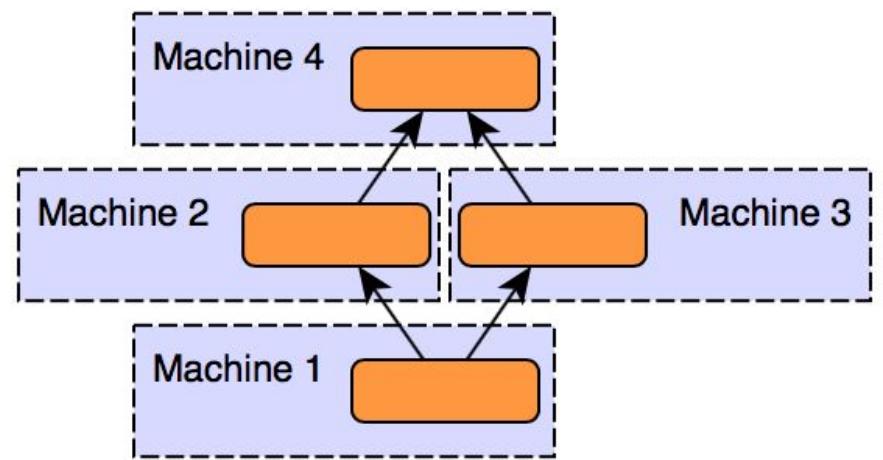
- They can even be used outside of Keras workflows!



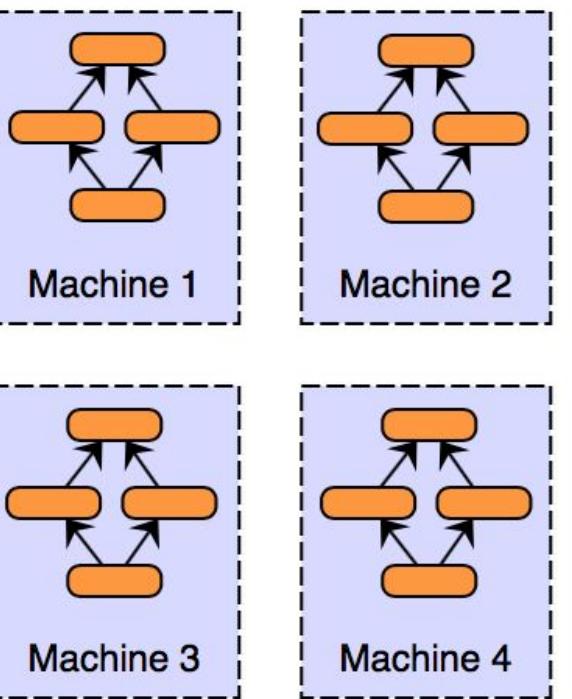
# Unified Distribution API



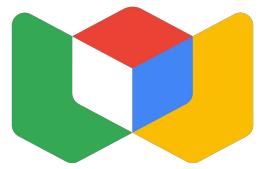
**Model Parallelism**



**Data Parallelism**



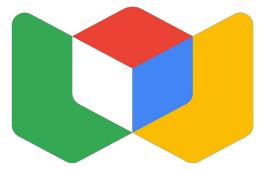
```
mesh_2d = keras.distribution.DeviceMesh(  
    shape=(2, 4), axis_names=["data", "model"], devices=devices  
)  
layout_map = keras.distribution.LayoutMap(mesh_2d)  
layout_map["d1/kernel"] = (None, "model")  
layout_map["d1/bias"] = ("model",)  
layout_map["d2/output"] = ("data", None)  
  
model_parallel = keras.distribution.ModelParallel(  
    mesh_2d, layout_map, batch_dim_name="data"  
)  
  
keras.distribution.set_distribution(model_parallel)  
  
inputs = layers.Input(shape=(28, 28, 1))  
y = layers.Flatten()(inputs)  
y = layers.Dense(units=200, use_bias=False, activation="relu",  
name="d1")(y)  
y = layers.Dropout(0.4)(y)  
y = layers.Dense(units=10, activation="softmax", name="d2")(y)  
model = keras.Model(inputs=inputs, outputs=y)  
  
model.compile(loss="mse")  
model.fit(dataset, epochs=3)  
model.evaluate(dataset)
```



# Keras Applications







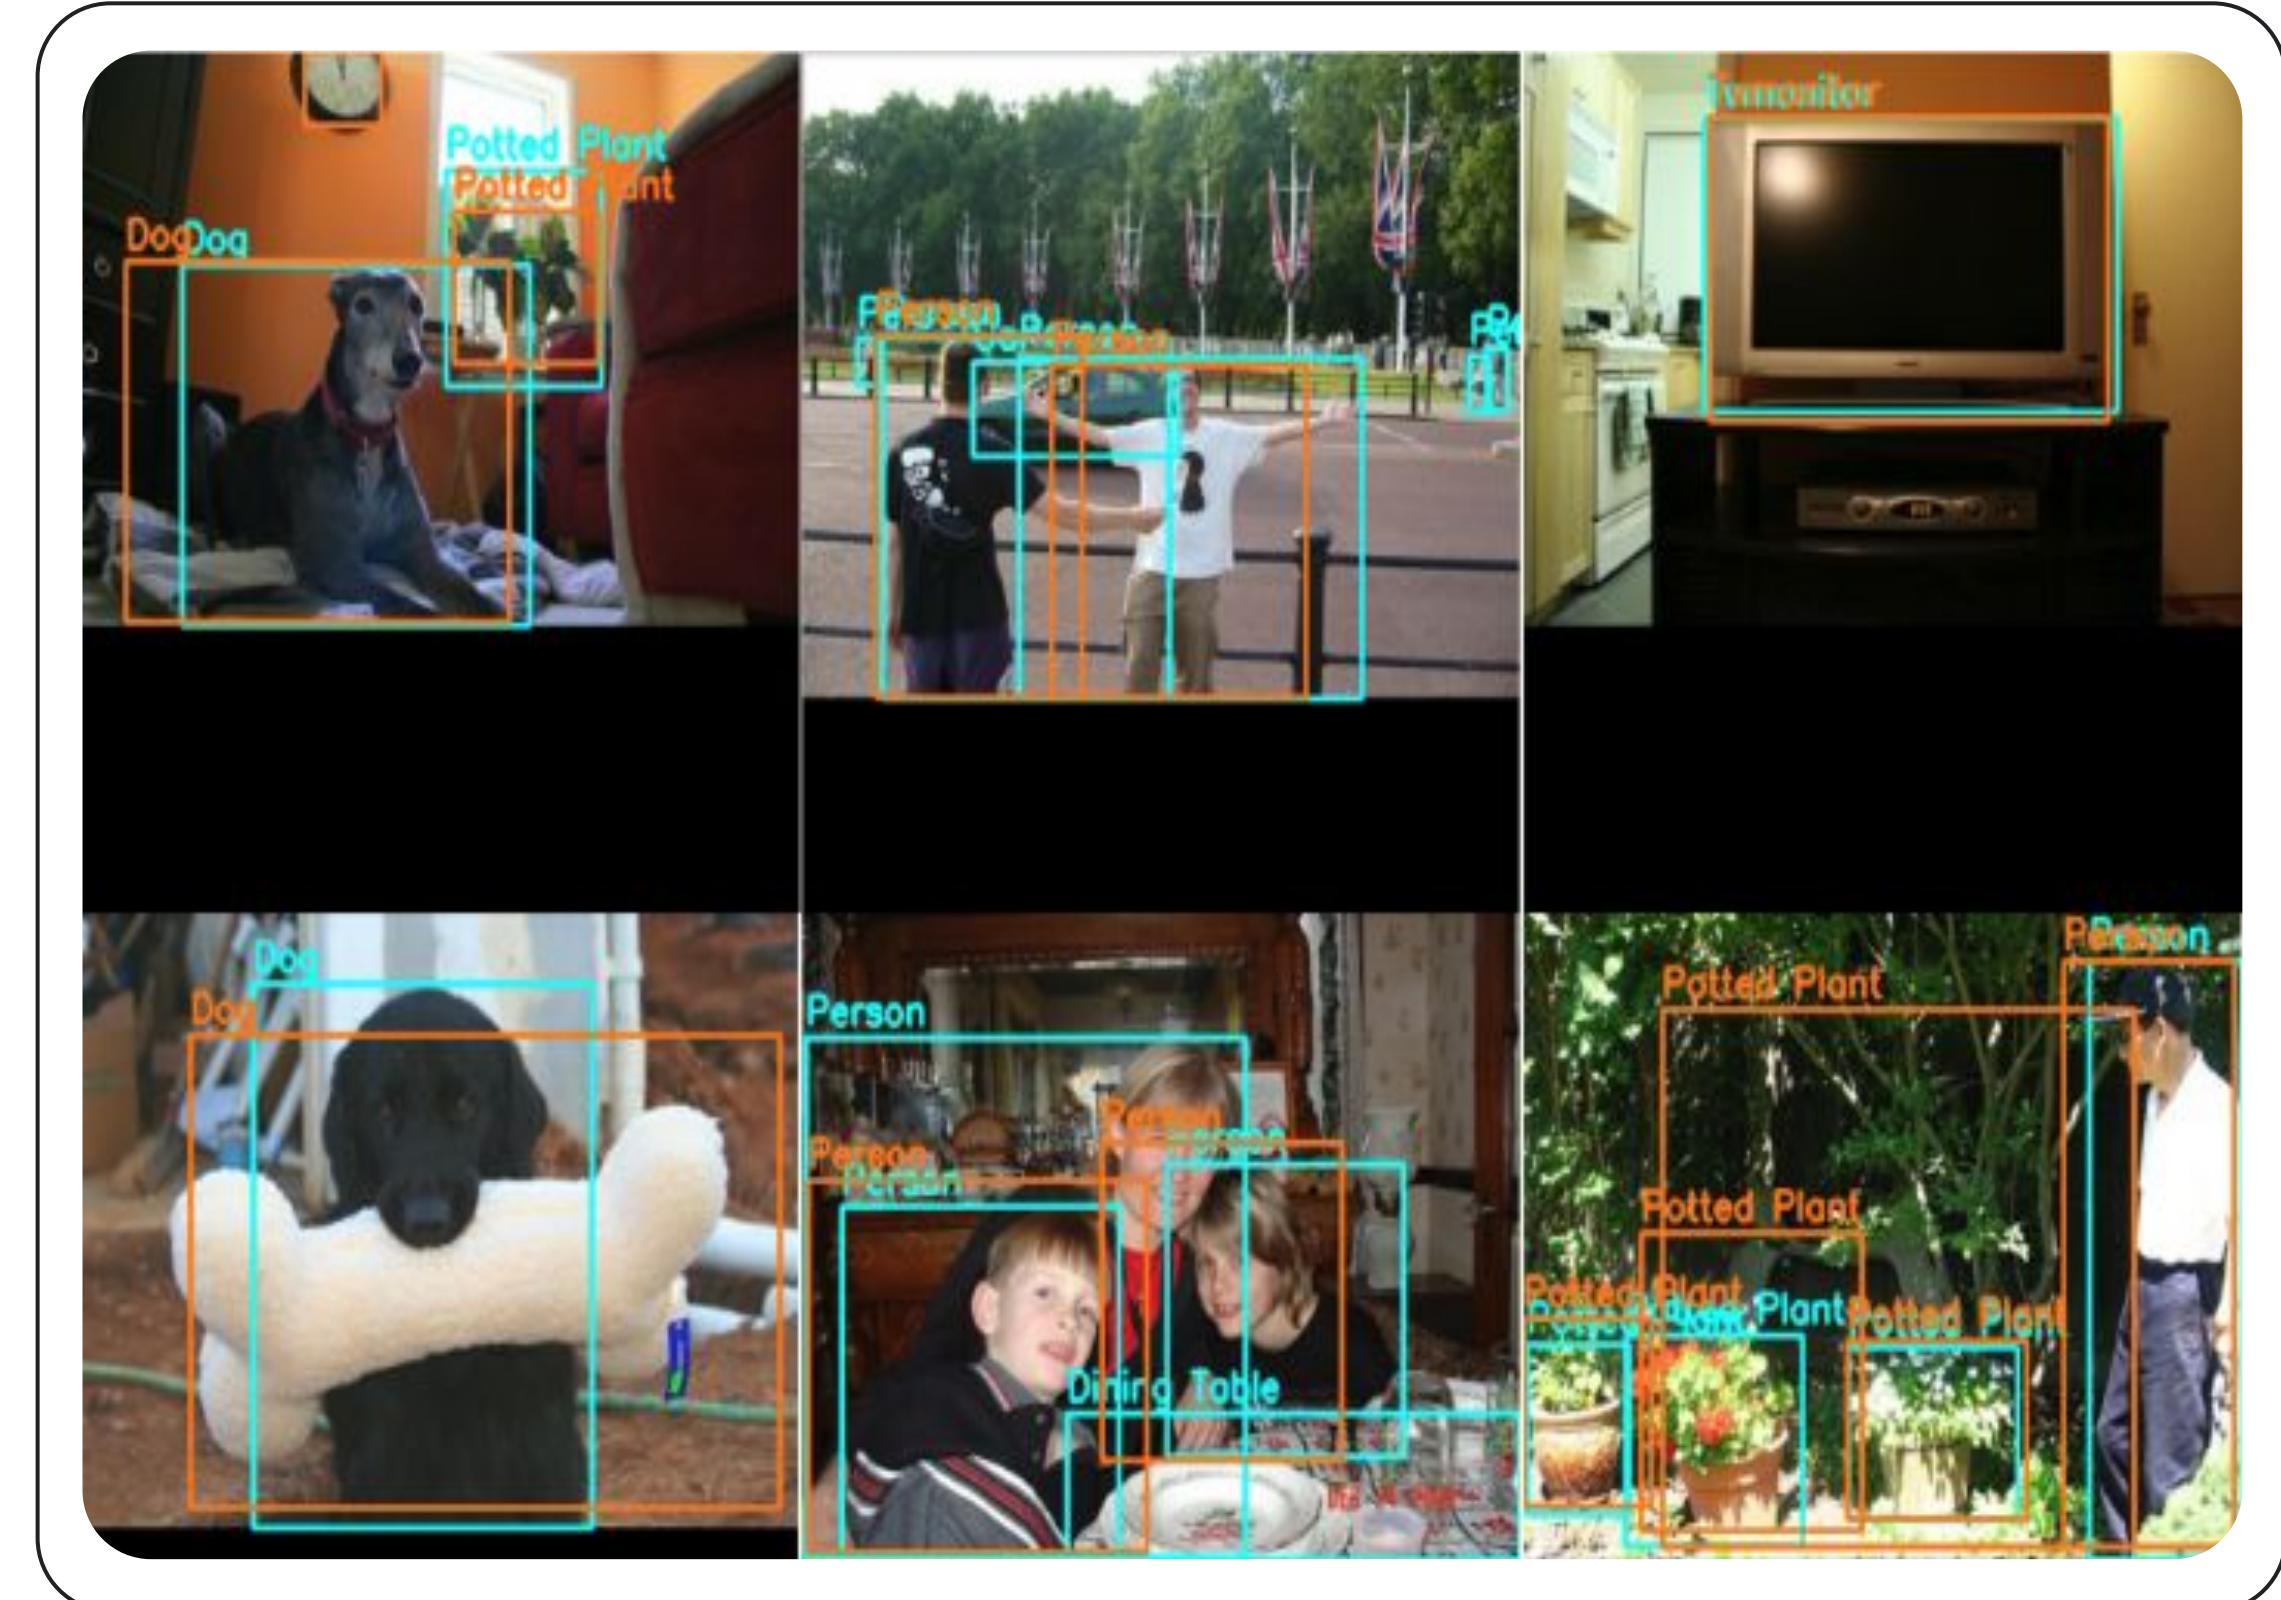
# Applied ML with KerasCV



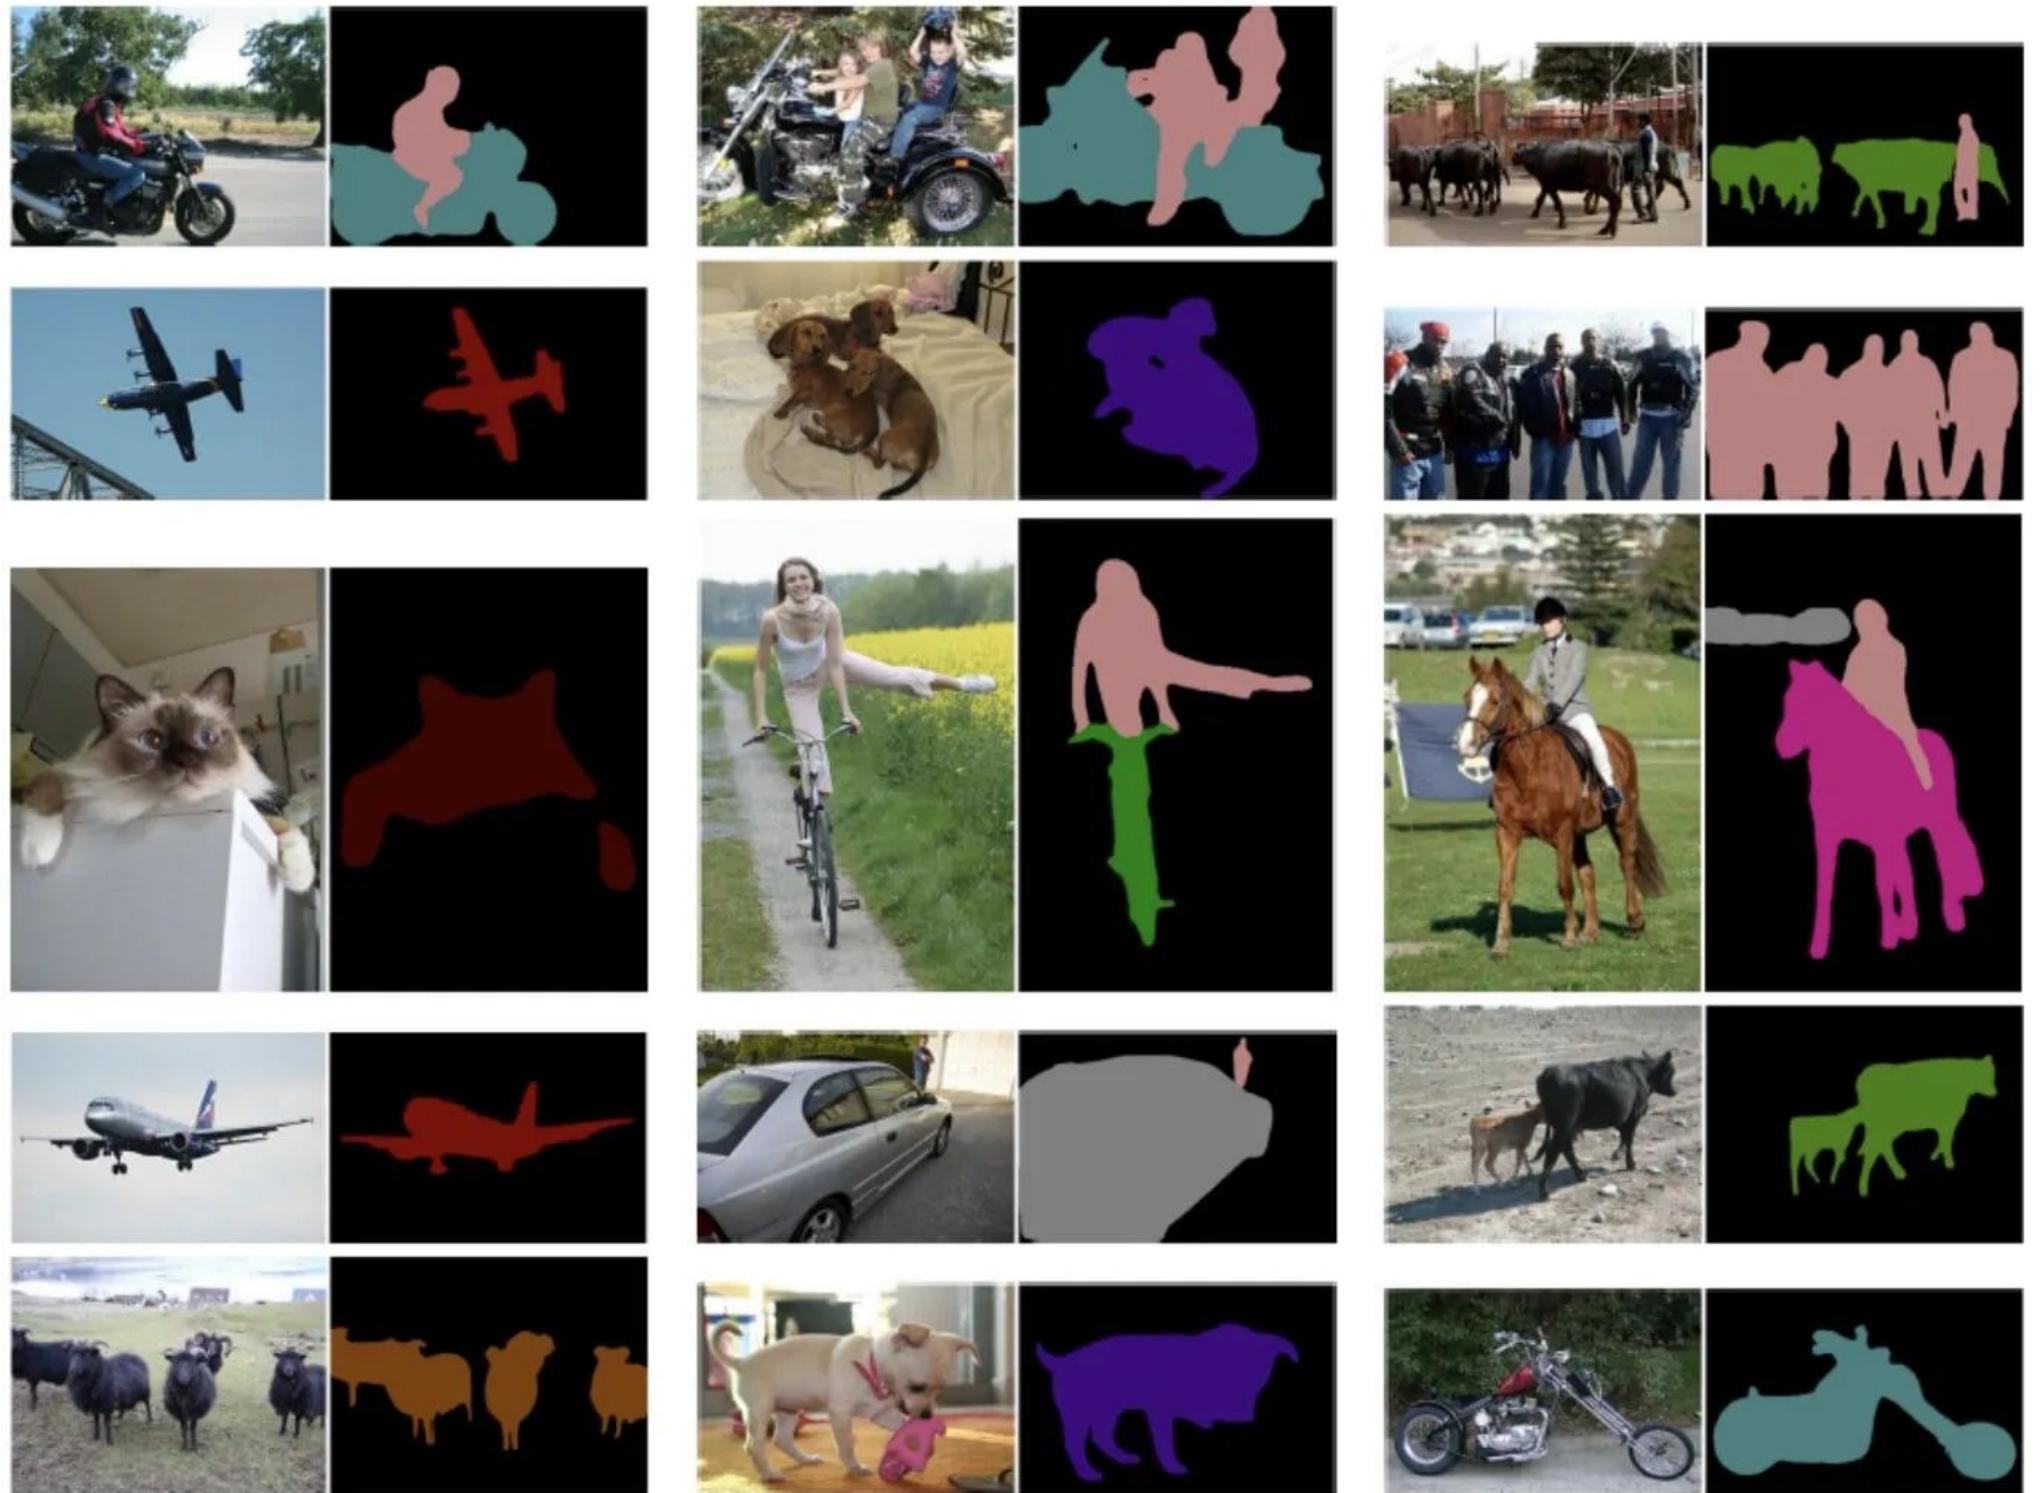
# Image Classification



# Object Detection

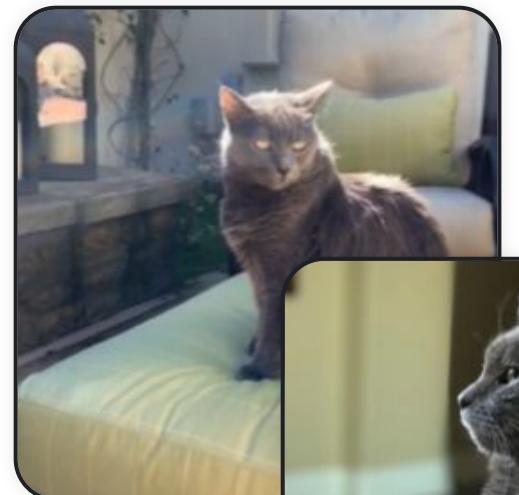


# Semantic Segmentation



# Image Generation





Textual inversion



As a fantasy character

Prompt to prompt

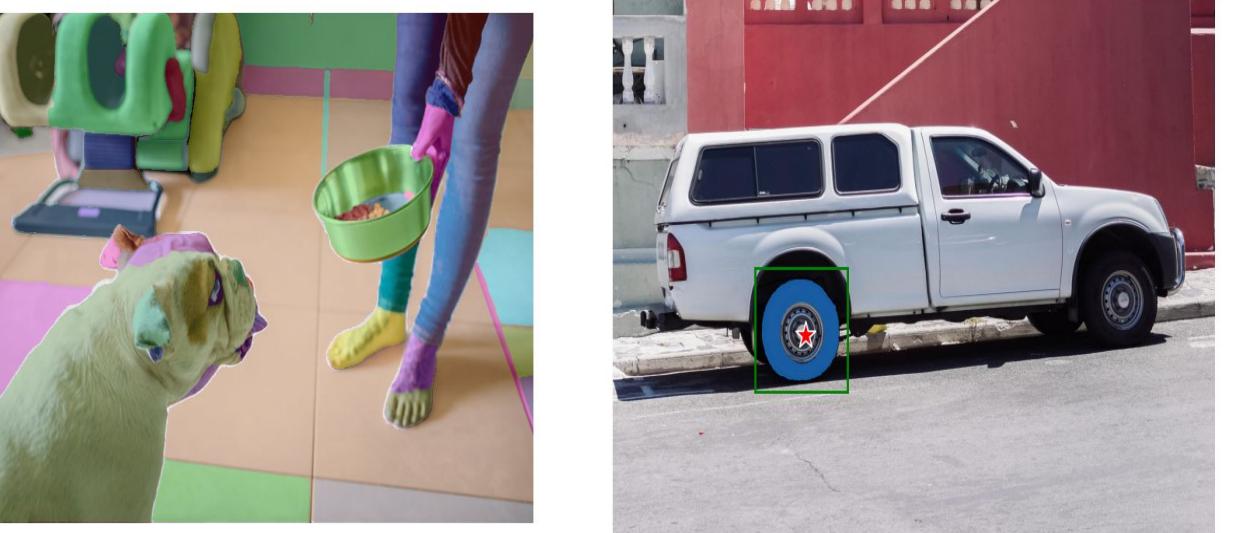


cat

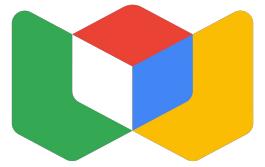
A photo of a **dog** with sunglasses



# Text prompted Segment Anything







# Applied ML with KerasNLP



# Text Classification

```
# Text classification
from keras_nlp.models import BertClassifier

model = BertClassifier.from_preset(
    "bert_base_en_uncased",
    num_classes=2,
    activation="softmax",
)

model.compile(...)

model.fit(imdb_movie_review_dataset)
model.predict(
    [
        "What an amazing movie!",
        "A total waste of my time.",
    ],
)
>>> array([[0.004173, 0.9956],
           [0.997, 0.0028]], dtype=float16)
```

# Text Generation

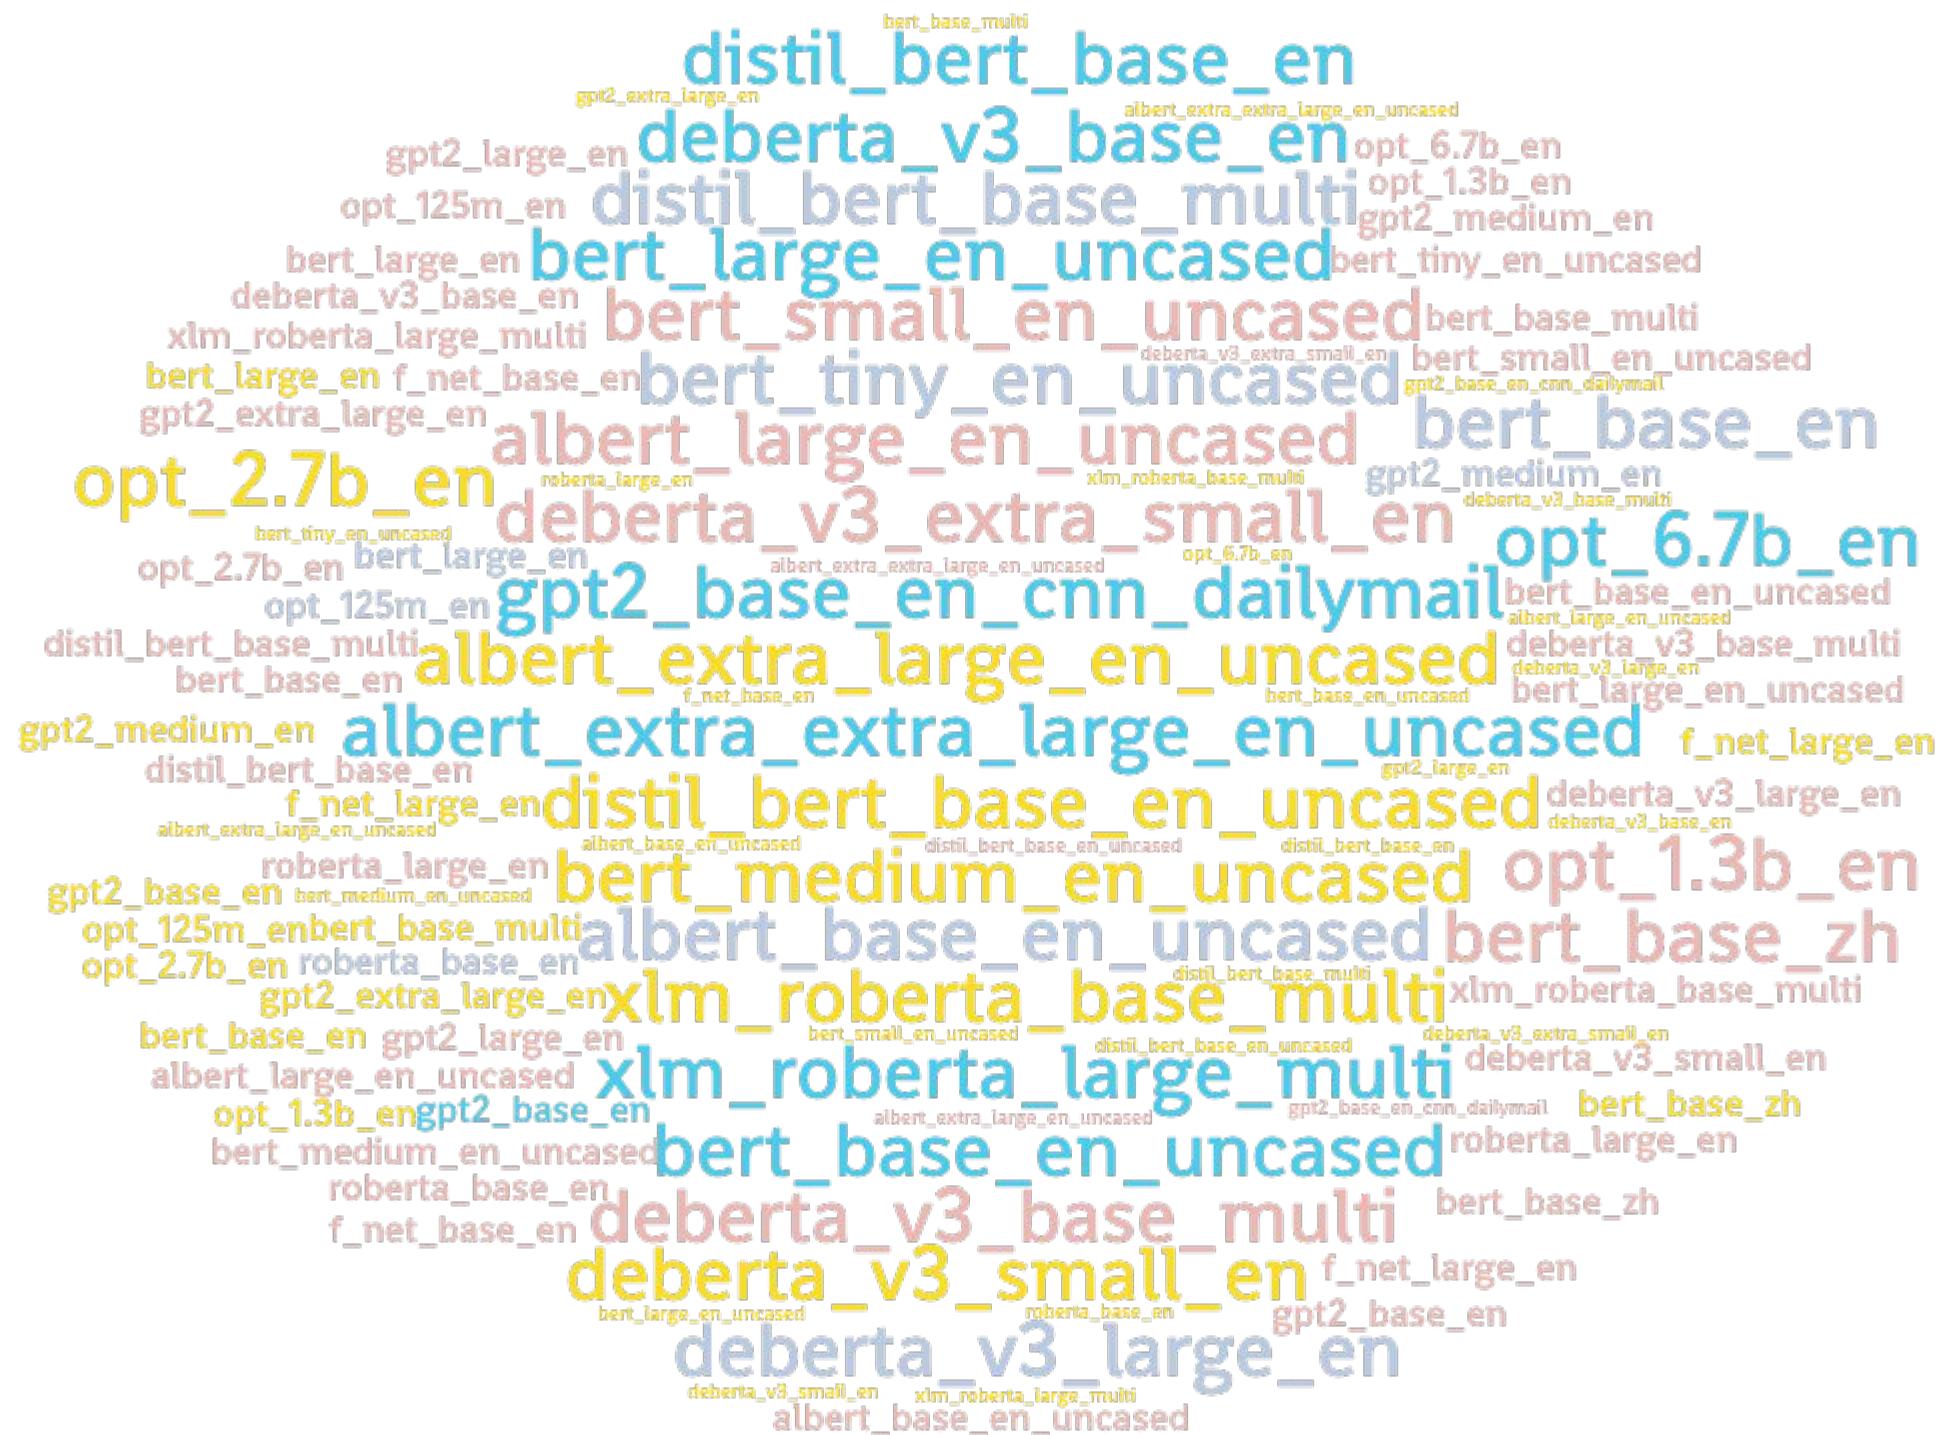
```
# Text generation
from keras_nlp.models import GPT2CausalLM

model = GPT2CausalLM.from_preset(
    "gpt2_base_en",
)

model.compile(...)

model.fit(cnn_dailymail_dataset)
model.generate(
    "Snowfall in Buffalo",
    max_length=40,
)

>>> 'Snowfall in Buffalo, New York, was expected to reach '
'2 feet by the end of the day, according to the National '
'Weather Service.'
```



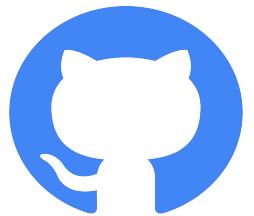
# Resources



Keras.io

- [https://keras.io/keras\\_core/](https://keras.io/keras_core/)
- [https://keras.io/keras\\_nlp/](https://keras.io/keras_nlp/)
- [https://keras.io/keras\\_cv/](https://keras.io/keras_cv/)

Github:



- <https://github.com/keras-team/keras-core>
- <https://github.com/keras-team/keras-nlp>
- <https://github.com/keras-team/keras-cv>



# Women in ML

# Thank you!



Divyashree Sreepathihalli (she/her)  
Software Engineer, Google  
@divyashreess

